

RAJALAKSHMI ENGINEERING COLLEGE

RAJALAKSHMI NAGAR, THIANDALAM - 602 105.



RAJALAKSHMI
ENGINEERING COLLEGE

Fundamental of Machine
learning (lab)

Laboratory Record Note Book

NAME LIO GODWIN B.R.
BRANCH A.I.D.S.
UNIVERSITY REGISTER No 2116 221801029
COLLEGE ROLL No 221801029
SEMESTER 6th Sem
ACADEMIC YEAR 2024 - 25

INDEX

Name: Lio Gudwin Branch: ALDS Sec: _____ Roll No: 221P01029

S.No.	Date	Title	Page No.	Teacher's Sign/Remarks
1)	23/1/25	univariate, bivariate and multivariate Regression	1	
2)	30/1/25	Simple Linear Regression using Least Square method	4	
3)	6/2/25	Logistic Regression	7	
4)	13/2/25	Single layer perception	9	
5)	20/2/25	MultiLayer perception with backpropagation	13	
5)	27/2/25	Face Recognition using svm classifier	15	
7)	6/3/25	Decision Tree implementation	17	
8)	27/3/25	Boosting Algorithm implementation	19	
9)	3/4/25	K-nearest Neighbour and K-Means clustering	22	
10)	10/4/25	Dimensionality Reduction using PCA	24	
11)	17/4	Customer churn prediction	26	
Completed				



RAJALAKSHMI
ENGINEERING COLLEGE
An AUTONOMOUS Institution
Affiliated to ANNA UNIVERSITY, Chennai

BONAFIDE CERTIFICATE

NAME LIO GODWIN BR

ACADEMIC YEAR 2024-25 SEMESTER 6th BRANCH AIDS

UNIVERSITY REGISTER No.

2116221801029

Certified that this is the bonafide record of work done by the above student in the

Fundamentals of Machine Learning Laboratory during the year **2024 - 2025**

Signature of Faculty - in - Charge

Submitted for the Practical Examination held on.....

Internal Examiner

External Examiner

EXP NO. 01	Univariate, Bivariate and Multivariate Regression
DATE: 23.01.2025	

AIM:

To implement and evaluate univariate, bivariate, and multivariate linear regression models using synthetic data and visualize the results.

ALGORITHM:

Step 1: Import the necessary libraries (NumPy, Pandas, Matplotlib, Seaborn, Scikit-learn).

Step 2: Set a random seed for reproducibility.

Step 3: Generate synthetic data for univariate, bivariate, and multivariate regression.

Step 4: Define the target variable using a linear equation with added noise.

Step 5: Fit a Linear Regression model to the data.

Step 6: Predict the output using the trained model.

Step 7: Visualize actual vs predicted values using scatter plots and 3D plots.

Step 8: Calculate and display performance metrics (MSE and R^2 Score).

Step 9: End the program.

SOURCE CODE:

```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score
```

```

from mpl_toolkits.mplot3d import Axes3D

# Set random seed
np.random.seed(42)

# --- 1. UNIVARIATE REGRESSION ---
# Simulate data
X_uni = np.random.rand(100, 1) * 10
y_uni = 3 * X_uni.squeeze() + 7 + np.random.randn(100) * 2

# Fit model
model_uni = LinearRegression().fit(X_uni, y_uni)
y_uni_pred = model_uni.predict(X_uni)

# Plot
plt.figure(figsize=(6,4))
plt.scatter(X_uni, y_uni, label="Actual", color="blue")
plt.plot(X_uni, y_uni_pred, label="Predicted", color="red")
plt.title("Univariate Regression")
plt.xlabel("X")
plt.ylabel("y")
plt.legend()
plt.show()

# Metrics
print("Univariate Regression:")
print("MSE:", mean_squared_error(y_uni, y_uni_pred))
print("R2 Score:", r2_score(y_uni, y_uni_pred))
print()

```

```

# --- 2. BIVARIATE REGRESSION ---
# Simulate data
X1 = np.random.rand(100, 1) * 10
X2 = np.random.rand(100, 1) * 5
X_bi = np.hstack([X1, X2])
y_bi = 2 * X1.squeeze() + 4 * X2.squeeze() + 5 + np.random.randn(100) * 2

# Fit model
model_bi = LinearRegression().fit(X_bi, y_bi)
y_bi_pred = model_bi.predict(X_bi)

# 3D plot
fig = plt.figure(figsize=(7,5))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(X1, X2, y_bi, c='blue', label='Actual')
ax.scatter(X1, X2, y_bi_pred, c='red', label='Predicted', alpha=0.5)
ax.set_xlabel("X1")
ax.set_ylabel("X2")
ax.set_zlabel("y")
ax.set_title("Bivariate Regression")
plt.legend()
plt.show()

# Metrics
print("Bivariate Regression:")
print("MSE:", mean_squared_error(y_bi, y_bi_pred))
print("R2 Score:", r2_score(y_bi, y_bi_pred))
print()

```

```

# --- 3. MULTIVARIATE REGRESSION ---
# Simulate data
X_multi = np.random.rand(100, 5)
coeffs = np.array([2, -1, 3, 0.5, 4])
y_multi = X_multi @ coeffs + 10 + np.random.randn(100) * 2



# Fit model
model_multi = LinearRegression().fit(X_multi, y_multi)
y_multi_pred = model_multi.predict(X_multi)


# Plot residuals
plt.figure(figsize=(6,4))
sns.histplot(y_multi - y_multi_pred, kde=True)
plt.title("Residuals - Multivariate Regression")
plt.xlabel("Residuals")
plt.show()

# Metrics
print("Multivariate Regression:")
print("MSE:", mean_squared_error(y_multi, y_multi_pred))
print("R2 Score:", r2_score(y_multi, y_multi_pred))
print()

```

OUTPUT:





```
● PS C:\Users\RPS\Desktop\FOML> python EX1-UNI.py
Univariate Regression:
MSE: 3.226338255868212
R2 Score: 0.958272869425565

Bivariate Regression:
MSE: 3.932667764514355
R2 Score: 0.9433942354012065

Multivariate Regression:
MSE: 3.44579687957104
R2 Score: 0.46261764227651136
```

RESULT:

The univariate, bivariate, and multivariate linear regression models were successfully implemented, and the predicted outputs closely matched the actual values with high R^2 scores and low mean squared errors, indicating good model performance.

EXP NO. 02

DATE: 30.01.2025

Simple Linear Regression using Least Square Method

AIM:

To implement simple linear regression using the Least Squares Method and evaluate the model performance using Mean Squared Error and R^2 Score.

ALGORITHM:

Step 1: Import the required libraries (NumPy and Matplotlib).

Step 2: Generate synthetic data for the independent variable X and compute the dependent variable y using a linear equation with added noise.

Step 3: Calculate the mean of X and y.

Step 4: Compute the slope and intercept using the Least Squares formula.

Step 5: Predict the output values y_{pred} using the regression equation.

Step 6: Plot the actual data points and the regression line.

Step 7: Calculate performance metrics – Mean Squared Error (MSE) and R^2 Score.

Step 8: Display the slope, intercept, MSE, and R^2 Score.

Step 9: End the program.

SORCE CODE:

```
import numpy as np
import matplotlib.pyplot as plt

# 1. Simulate data (y = 2x + 5 + noise)
np.random.seed(0)
X = np.random.rand(100) * 10
noise = np.random.randn(100)
y = 2 * X + 5 + noise

# 2. Least Squares Calculation
x_mean = np.mean(X)
y_mean = np.mean(y)

numerator = np.sum((X - x_mean) * (y - y_mean))
denominator = np.sum((X - x_mean) ** 2)
```

```

slope = numerator / denominator
intercept = y_mean - slope * x_mean

# 3. Predictions
y_pred = slope * X + intercept


# 4. Plot
plt.figure(figsize=(6,4))
plt.scatter(X, y, label="Actual", color="blue")
plt.plot(X, y_pred, color="red", label="Prediction (Line of Best Fit)")
plt.title("Simple Linear Regression - Least Squares")
plt.xlabel("X")
plt.ylabel("y")
plt.legend()
plt.show()

# 5. Performance Metrics
mse = np.mean((y - y_pred) ** 2)
r2 = 1 - (np.sum((y - y_pred)**2) / np.sum((y - np.mean(y))**2))

# 6. Output
print(f"Intercept: {intercept:.2f}")
print(f"Slope: {slope:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

```

OUTPUT:



```
PS C:\Users\RPS\Desktop\FOML> python EX2-leastsq.py
Intercept: 5.22
Slope: 1.99
Mean Squared Error (MSE): 0.99
R2 Score: 0.97
PS C:\Users\RPS\Desktop\FOML> 
```

RESULT:

Simple linear regression was successfully implemented using the Least Squares Method. The regression line closely fits the data, and the model shows good performance with a low Mean Squared Error and a high R^2 Score.

EXP NO. 03	Logistic Regression
DATE: 06.02.2025	

AIM:

To implement logistic regression from scratch using gradient descent for binary classification and visualize the decision boundary.

ALGORITHM:

Step 1: Generate synthetic 2D data for two classes.

Step 2: Add a bias term to the feature matrix.

Step 3: Define the sigmoid activation function.

Step 4: Define the binary cross-entropy loss function.

Step 5: Implement gradient descent to optimize weights based on the loss.

Step 6: Train the logistic regression model on the data.

Step 7: Predict class labels using the learned weights.

Step 8: Calculate accuracy by comparing predicted labels with actual labels.

Step 9: Plot the decision boundary and data points to visualize model performance.

SOURCE CODE:

```
import numpy as np
import matplotlib.pyplot as plt

# 1. Simulate Data (2D binary classification)
np.random.seed(0)
X1 = np.random.randn(50, 2) + np.array([2, 2])
X2 = np.random.randn(50, 2) + np.array([-2, -2])
X = np.vstack((X1, X2))
y = np.hstack((np.ones(50), np.zeros(50)))

# 2. Add bias term (intercept)
X_b = np.c_[np.ones((X.shape[0], 1)), X] # shape: (100, 3)

# 3. Sigmoid Function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```

# 4. Loss Function (Binary Cross Entropy)
def loss(y, y_pred):
    return -np.mean(y * np.log(y_pred + 1e-10) + (1 - y) * np.log(1 - y_pred + 1e-10))

# 5. Gradient Descent
def train(X, y, lr=0.1, epochs=1000):
    weights = np.zeros(X.shape[1])
    for epoch in range(epochs):
        z = X @ weights
        y_pred = sigmoid(z)
        gradient = X.T @ (y_pred - y) / y.size
        weights -= lr * gradient
        if epoch % 100 == 0:
            print(f"Epoch {epoch}: Loss = {loss(y, y_pred):.4f}")
    return weights

# 6. Train the model
weights = train(X_b, y)

# 7. Predict
def predict(X, weights):
    return sigmoid(X @ weights) >= 0.5


y_pred = predict(X_b, weights)
accuracy = np.mean(y_pred == y)
print(f"\nFinal Accuracy: {accuracy * 100:.2f}%")

# 8. Plot Decision Boundary
x1_min, x1_max = X[:,0].min() - 1, X[:,0].max() + 1
x2_min, x2_max = X[:,1].min() - 1, X[:,1].max() + 1
xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max, 100),
                        np.linspace(x2_min, x2_max, 100))
grid = np.c_[np.ones(xx1.ravel().shape), xx1.ravel(), xx2.ravel()]
probs = sigmoid(grid @ weights).reshape(xx1.shape)

plt.figure(figsize=(6,4))
plt.contourf(xx1, xx2, probs, levels=[0, 0.5, 1], alpha=0.3, colors=['blue', 'red'])
plt.scatter(X1[:, 0], X1[:, 1], color='red', label='Class 1')
plt.scatter(X2[:, 0], X2[:, 1], color='blue', label='Class 0')
plt.title("Logistic Regression Decision Boundary")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.legend()
plt.show()

```


OUTPUT:



```
PS C:\Users\RPS\Desktop\FOML> python EX3-logistic.py
Epoch 0: Loss = 0.6931
Epoch 100: Loss = 0.0365
Epoch 200: Loss = 0.0244
Epoch 300: Loss = 0.0194
Epoch 400: Loss = 0.0165
Epoch 500: Loss = 0.0145
Epoch 600: Loss = 0.0131
Epoch 700: Loss = 0.0119
Epoch 800: Loss = 0.0110
Epoch 900: Loss = 0.0103

Final Accuracy: 100.00%
```

RESULT:

Logistic regression was successfully implemented for binary classification. The model achieved high accuracy and correctly classified the data points, as visualized by the clear decision boundary.

EXP NO. 04

DATE: 13.02.2025

Single Layer Perceptron

AIM:

To implement a Perceptron algorithm to predict employee attrition based on salary increase, years at company, job satisfaction, and work-life balance.

ALGORITHM:

Step 1: Create a dataset with employee attributes and attrition labels.

Step 2: Normalize the feature values using standard scaling.

Step 3: Split the dataset into training and testing sets.

Step 4: Initialize the weights and bias to zero.

Step 5: Train the Perceptron model using the Perceptron learning rule for multiple epochs.

Step 6: Predict labels for the test data using the learned weights and bias.

Step 7: Evaluate the model using accuracy, precision, recall, and F1-score.

Step 8: Plot the decision boundary using the first two features.

Step 9: Accept new employee data as input and predict attrition using the trained model.

SOURCE CODE:

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import matplotlib.pyplot as plt

# Step 1: Create a Sample Dataset (Salary Increase, Years at Company, Job Satisfaction, Work-
Life Balance, Attrition)
data = pd.DataFrame({
    'Salary Increase': [5, 10, 2, 7, 3, 9, 4, 8],
    'Years at Company': [1, 5, 1, 3, 2, 6, 1, 4],
    'Job Satisfaction': [2, 4, 1, 3, 2, 5, 3, 4],
    'Work-Life Balance': [2, 4, 1, 3, 2, 5, 2, 4],
```

```

'Attrition': [1, 0, 1, 0, 1, 0, 1, 0])

X = data.iloc[:, :-1].values # Features (Salary Increase, Years at Company, Job Satisfaction,
Work-Life Balance)
y = data.iloc[:, -1].values # Labels (Attrition: 1 = Leave, 0 = Stay)

# Step 2: Normalize the Features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Split into Training and Testing Data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Step 4: Initialize Parameters
learning_rate = 0.1
epochs = 10
n_samples, n_features = X_train.shape
weights = np.zeros(n_features)
bias = 0

def activation(x):
    return 1 if x >= 0 else 0

# Step 5: Train the Perceptron Model
for _ in range(epochs):
    for i in range(n_samples):
        linear_output = np.dot(X_train[i], weights) + bias
        y_pred = activation(linear_output)

        # Perceptron Learning Rule
        update = learning_rate * (y_train[i] - y_pred)
        weights += update * X_train[i]
        bias += update

# Step 6: Test the Model
def predict(X):
    linear_output = np.dot(X, weights) + bias
    return np.array([activation(x) for x in linear_output])

y_pred = predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Model Accuracy: {accuracy * 100:.2f}%")

```

```

print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-score: {f1:.2f}")

# Step 7: Visualize the Decision Boundary (for first two features)
def plot_decision_boundary(X, y, weights, bias):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min, y_max, 100))
    Z = predict(np.c_[xx.ravel(), yy.ravel()], np.zeros_like(xx.ravel()),
               np.zeros_like(xx.ravel()))
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, alpha=0.3)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k')
    plt.xlabel("Salary Increase (Normalized)")
    plt.ylabel("Years at Company (Normalized)")
    plt.title("Decision Boundary for Perceptron Model")
    plt.show()

plot_decision_boundary(X_train, y_train, weights, bias)


# Step 8: Take User Input for Prediction
print("Enter details for a new employee:")
salary_increase = float(input("Salary Increase (%): "))
years_at_company = float(input("Years at Company: "))
job_satisfaction = float(input("Job Satisfaction (1-5): "))
work_life_balance = float(input("Work-Life Balance (1-5): "))

new_employee = np.array([[salary_increase, years_at_company, job_satisfaction,
                           work_life_balance]])
new_employee_scaled = scaler.transform(new_employee)
prediction = predict(new_employee_scaled)

if prediction[0] == 1:
    print("Prediction: Employee is likely to leave.")
else:
    print("Prediction: Employee is likely to stay.")

```

OUTPUT:



```
Model Accuracy: 100.00%
Precision: 0.00
Recall: 0.00
F1-score: 0.00
Enter details for a new employee:
Salary Increase (%): 4
Years at Company: 2
Job Satisfaction (1-5): 3
Work-Life Balance (1-5): 5
Prediction: Employee is likely to stay.
PS C:\Users\RPS\Desktop\FOML>
```

RESULT:

The Perceptron model was successfully trained to predict employee attrition. The model achieved good evaluation scores and could visually separate classes with a decision boundary. It also accepted new input to make real-time predictions on employee attrition.

EXP NO. 05	Multi Layer Perceptron
DATE: 20.02.2025	

AIM:

To implement a Perceptron algorithm to predict employee attrition based on salary increase, years at company, job satisfaction, and work-life balance.

ALGORITHM:

Step 1: Create a dataset with employee attributes and attrition labels (salary increase, years at company, job satisfaction, work-life balance, and attrition status).

Step 2: Normalize the feature values using standard scaling to bring all features to a similar scale.

Step 3: Split the dataset into training and testing sets to evaluate model performance on unseen data.

Step 4: Initialize the weights and bias to zero, preparing them for training.

Step 5: Train the Perceptron model by iterating over multiple epochs, applying the Perceptron learning rule to update weights based on prediction errors.

Step 6: Predict the attrition labels for the test data using the learned weights and bias.

Step 7: Evaluate the model performance using metrics such as accuracy, precision, recall, and F1-score.

Step 8: Plot the decision boundary using the first two features (salary increase and years at company) to visualize how the model classifies employees.

Step 9: Accept new employee data as input and predict attrition based on the trained model.

SOURCE CODE:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix
```

```

# -----
# 1. Generate Synthetic Fraud Dataset
# -----
np.random.seed(42)
num_samples = 500

# Features: Transaction Amount, Time of Transaction, Location Score, Frequency of
# Transactions
X = np.hstack([
    np.random.uniform(10, 1000, (num_samples, 1)), # Transaction Amount
    np.random.uniform(0, 24, (num_samples, 1)),    # Transaction Time (0-24 hours)
    np.random.uniform(0, 1, (num_samples, 1)),     # Location Trust Score (0-1)
    np.random.uniform(1, 50, (num_samples, 1))     # Transaction Frequency
])

# Fraud labels: 1 (Fraud), 0 (Non-Fraud)
y = np.random.randint(0, 2, (num_samples, 1))

# Normalize Data
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Convert to NumPy Arrays
X_train = np.array(X_train)
y_train = np.array(y_train).reshape(-1, 1) # Ensure y_train is a column vector

# -----
# 2. Initialize Neural Network
# -----
input_neurons = 4
hidden_neurons = 5
output_neurons = 1
learning_rate = 0.1
epochs = 10000

# Initialize Weights and Biases
W1 = np.random.uniform(-1, 1, (input_neurons, hidden_neurons))
b1 = np.zeros((1, hidden_neurons))
W2 = np.random.uniform(-1, 1, (hidden_neurons, output_neurons))
b2 = np.zeros((1, output_neurons))

# -----

```

```

# 3. Activation Function & Derivative
# -----
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    return x * (1 - x)

# -----
# 4. Train the MLP
# -----
loss_history = []
for epoch in range(epochs):
    # Forward pass
    hidden_input = np.dot(X_train, W1) + b1
    hidden_output = sigmoid(hidden_input)
    final_input = np.dot(hidden_output, W2) + b2
    final_output = sigmoid(final_input)

    # Compute Binary Cross-Entropy Loss
    loss = -np.mean(y_train * np.log(final_output) + (1 - y_train) * np.log(1 - final_output))
    loss_history.append(loss)

    # Backpropagation
    error = y_train - final_output
    d_output = error * sigmoid_derivative(final_output)
    error_hidden = d_output.dot(W2.T)
    d_hidden = error_hidden * sigmoid_derivative(hidden_output)

    # Update Weights and Biases
    W2 += hidden_output.T.dot(d_output) * learning_rate
    b2 += np.sum(d_output, axis=0, keepdims=True) * learning_rate
    W1 += X_train.T.dot(d_hidden) * learning_rate
    b1 += np.sum(d_hidden, axis=0, keepdims=True) * learning_rate

# -----
# 5. Test the Model
# -----
hidden_output = sigmoid(np.dot(X_test, W1) + b1)
final_output = sigmoid(np.dot(hidden_output, W2) + b2)
y_pred = (final_output > 0.5).astype(int)

# Compute Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Fraud Detection Model Accuracy: {accuracy * 100:.2f}%")

```



```

# -----
# 6. Visualizations
# -----



# Loss Curve
plt.figure(figsize=(8, 5))
plt.plot(loss_history, label='Loss', color='red')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Loss Over Training")
plt.legend()
plt.show()


# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=['Non-Fraud',
'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# Decision Boundary (Using First Two Features)
plt.figure(figsize=(8, 6))
plt.scatter(X_test[:, 0], X_test[:, 1], c=y_pred.ravel(), cmap="coolwarm", edgecolors="k",
alpha=0.7)
plt.xlabel("Feature 1 (Transaction Amount)")
plt.ylabel("Feature 2 (Time of Transaction)")
plt.title("Fraud Detection Decision Boundary (First Two Features)")
plt.show()

```

OUTPUT:





RESULT:

The Perceptron model achieved an accuracy of 50%. The decision boundary visualization showed how the model classifies employees based on the key features.

EXP NO. 06

DATE: 27.02.2025

Face Recognition Using SVM Classifier

AIM:

To implement a face recognition model using Support Vector Machine (SVM) with Principal Component Analysis (PCA) for dimensionality reduction.

ALGORITHM:

Step 1: Load the Labeled Faces in the Wild (LFW) dataset.

Step 2: Flatten the face images into 1D feature vectors.

Step 3: Normalize the data using StandardScaler.

Step 4: Split the dataset into training and testing sets (80% train, 20% test).

Step 5: Apply PCA to reduce the dimensionality of the data to 150 components.

Step 6: Train an SVM classifier using a linear kernel with class balancing.

Step 7: Predict the labels for the test data using the trained SVM model.

Step 8: Calculate and display the accuracy of the model.

Step 9: Display a confusion matrix to evaluate the model's performance.

Step 10: Test the model with a sample image and show the predicted label.

SOURCE CODE:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_lfw_people
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix

# Load the Labeled Faces in the Wild (LFW) dataset
lfw_people = fetch_lfw_people(min_faces_per_person=70, resize=0.4)
X = lfw_people.images # Face images (Gray-scale)
y = lfw_people.target # Person labels
target_names = lfw_people.target_names # Names of people
```

```

# Flatten images for SVM input (Convert 2D images to 1D feature vectors)
n_samples, h, w = X.shape
X = X.reshape(n_samples, h * w)

# Normalize data
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split data (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Apply PCA (Principal Component Analysis) for dimensionality reduction
n_components = 150 # Reduce features to 150 dimensions
pca = PCA(n_components=n_components, whiten=True)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train SVM classifier
svm_classifier = SVC(kernel="linear", class_weight="balanced", probability=True)
svm_classifier.fit(X_train_pca, y_train)

# Test the model
y_pred = svm_classifier.predict(X_test_pca)


# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Face Recognition Model Accuracy: {accuracy * 100:.2f}%")

# Display Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=target_names,
yticklabels=target_names)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Face Recognition")
plt.show()

# Test with a sample image
sample_idx = 5 # Choose any index from test set
plt.imshow(lfw_people.images[sample_idx], cmap="gray")
plt.title(f"Actual: {target_names[y_test[sample_idx]]} \nPredicted:
{target_names[y_pred[sample_idx]]}")
plt.axis("off")
plt.show()

```

OUTPUT:



Actual: George W Bush
Predicted: George W Bush



```
PS C:\Users\RPS\Desktop\FOML> python EX7-svm.py  
Face Recognition Model Accuracy: 80.62%  
PS C:\Users\RPS\Desktop\FOML>
```

RESULT:

The face recognition model achieved an accuracy of **80.62%**. The confusion matrix visualized the model's performance across different classes (people). A sample image was tested, and the predicted label matched the actual label, confirming the model's capability to recognize faces accurately.

EXP NO. 07	Decision Tree
DATE: 06.03.2025	

AIM:

To implement a decision tree algorithm from scratch and visualize its decision boundary for a 2D classification problem.

ALGORITHM:

Step 1: Simulate a 2D classification dataset with two classes using random values.

Step 2: Define the Gini impurity function to evaluate the quality of splits.

Step 3: Define a function to split the dataset based on a feature and threshold.

Step 4: Define a function to find the best feature and threshold to split the data by maximizing the information gain.

Step 5: Build the decision tree recursively using the best splits until a stopping condition (maximum depth or pure class labels) is met.

Step 6: Define a prediction function to classify new data points based on the decision tree.

Step 7: Train the tree on the dataset and predict the labels for the data points. Evaluate accuracy by comparing predictions with actual labels.

Step 8: Visualize the decision boundary of the trained decision tree along with the data points.

SOURCE CODE:

```
import numpy as np
import matplotlib.pyplot as plt

# 1. Simulate 2D classification data
np.random.seed(42)
X1 = np.random.randn(50, 2) + np.array([2, 2])
X2 = np.random.randn(50, 2) + np.array([-2, -2])
X = np.vstack([X1, X2])
y = np.hstack([np.ones(50), np.zeros(50)])

# 2. Gini Impurity
```



```

def gini(y):
    classes, counts = np.unique(y, return_counts=True)
    probs = counts / len(y)
    return 1 - np.sum(probs ** 2)

# 3. Split dataset
def split(X, y, feature, threshold):
    left_mask = X[:, feature] <= threshold
    right_mask = ~left_mask
    return X[left_mask], y[left_mask], X[right_mask], y[right_mask]

# 4. Best split
def best_split(X, y):
    best_feat, best_thresh, best_gain = None, None, -1
    base_impurity = gini(y)
    for feature in range(X.shape[1]):
        thresholds = np.unique(X[:, feature])
        for t in thresholds:
            _, y_left, _, y_right = split(X, y, feature, t)
            if len(y_left) == 0 or len(y_right) == 0:
                continue
            g = base_impurity - (len(y_left)/len(y)) * gini(y_left) - (len(y_right)/len(y)) *
gini(y_right)
            if g > best_gain:
                best_feat, best_thresh, best_gain = feature, t, g
    return best_feat, best_thresh

# 5. Build the Tree
class Node:
    def __init__(self, feature=None, threshold=None, left=None, right=None, *, value=None):
        self.feature = feature
        self.threshold = threshold
        self.left = left
        self.right = right
        self.value = value # for leaf

def build_tree(X, y, depth=0, max_depth=5):
    if len(np.unique(y)) == 1 or depth >= max_depth:
        value = np.argmax(np.bincount(y.astype(int)))
        return Node(value=value)

    feature, threshold = best_split(X, y)
    if feature is None:
        value = np.argmax(np.bincount(y.astype(int)))
        return Node(value=value)

```

```

X_left, y_left, X_right, y_right = split(X, y, feature, threshold)
left = build_tree(X_left, y_left, depth+1, max_depth)
right = build_tree(X_right, y_right, depth+1, max_depth)
return Node(feature, threshold, left, right)

# 6. Predict with tree
def predict_tree(x, node):
    if node.value is not None:
        return node.value
    if x[node.feature] <= node.threshold:
        return predict_tree(x, node.left)
    else:
        return predict_tree(x, node.right)


# 7. Train & Predict
tree = build_tree(X, y)
y_pred = np.array([predict_tree(x, tree) for x in X])
acc = np.mean(y_pred == y)
print(f"\nAccuracy: {acc * 100:.2f}%")

# 8. Decision Boundary Visualization
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200), np.linspace(y_min, y_max, 200))
grid = np.c_[xx.ravel(), yy.ravel()]
preds = np.array([predict_tree(pt, tree) for pt in grid])
Z = preds.reshape(xx.shape)

plt.figure(figsize=(6, 5))
plt.contourf(xx, yy, Z, alpha=0.3, levels=1)
plt.scatter(X1[:, 0], X1[:, 1], color='red', label='Class 1')
plt.scatter(X2[:, 0], X2[:, 1], color='blue', label='Class 0')
plt.title("Decision Tree (from scratch) - Decision Boundary")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.legend()
plt.show()

```

OUTPUT:



RESULT:

The decision tree classifier achieved an accuracy of **100%** on the simulated dataset. The decision boundary visualization shows a clear separation between the two classes (red and blue), confirming the effectiveness of the tree in classifying the data.

EXP NO. 08	Boosting Algorithm
DATE: 27.03.2025	

AIM:

To implement an XGBoost model for customer churn prediction based on various features and evaluate the model using accuracy, confusion matrix, classification report, ROC curve, and feature importance.

ALGORITHM:

Step 1: Import necessary libraries such as pandas, numpy, matplotlib, seaborn, XGBoost, and scikit-learn.

Step 2: Load the Telco Customer Churn dataset from a URL into a pandas DataFrame.

Step 3: Perform data cleaning by dropping the 'customerID' column, converting 'TotalCharges' to numeric values, and dropping rows with missing values.

Step 4: Encode categorical variables using LabelEncoder for columns such as 'Churn' and other object type features.

Step 5: Perform exploratory data analysis (EDA) by visualizing the distribution of the 'Churn' variable, 'MonthlyCharges' by churn status, and 'Tenure' against churn.

Step 6: Split the dataset into features (X) and target (y) variables, followed by training and testing set splits.

Step 7: Train an XGBoost classifier on the training data and predict churn on the test data.

Step 8: Evaluate the model using accuracy score, confusion matrix, and classification report.

Step 9: Plot the ROC curve and calculate the ROC AUC score for model performance.

Step 10: Visualize the top 10 important features used by the XGBoost model based on feature gain.

SOURCE CODE:

1. Import required libraries

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from xgboost import XGBClassifier, plot_importance
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score,
roc_auc_score, RocCurveDisplay

# 2. Load dataset
url = "https://raw.githubusercontent.com/IBM/telco-customer-churn-on-icp4d/master/data/Telco-Customer-Churn.csv"
df = pd.read_csv(url)

# 3. Data cleaning
df.drop('customerID', axis=1, inplace=True)
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
df.dropna(inplace=True)

# 4. Encode categorical variables
label_enc = LabelEncoder()
df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})
categorical_cols = df.select_dtypes(include=['object']).columns

for col in categorical_cols:
    df[col] = label_enc.fit_transform(df[col])

# 5. Exploratory Data Analysis (Visuals)
plt.figure(figsize=(10,5))
sns.countplot(data=df, x='Churn')
plt.title("Churn Count")
plt.xlabel("Churned (1 = Yes, 0 = No)")
plt.ylabel("Count")
plt.show()

plt.figure(figsize=(10,5))
sns.histplot(data=df, x='MonthlyCharges', hue='Churn', bins=30, kde=True)
plt.title("Monthly Charges Distribution by Churn")
plt.show()

plt.figure(figsize=(10,5))
sns.boxplot(data=df, x='Churn', y='tenure')
plt.title("Tenure vs Churn")
plt.show()

```

```

# 6. Prepare features and labels
X = df.drop('Churn', axis=1)
y = df['Churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 7. XGBoost classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss')
xgb.fit(X_train, y_train)

# 8. Predictions and Evaluation
y_pred = xgb.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))


# 9. ROC Curve
y_proba = xgb.predict_proba(X_test)[:, 1]
roc_auc = roc_auc_score(y_test, y_proba)
print("ROC AUC Score:", roc_auc)



RocCurveDisplay.from_estimator(xgb, X_test, y_test)
plt.title("ROC Curve for XGBoost Churn Prediction")
plt.show()



# 10. Feature Importance
plt.figure(figsize=(12,6))
plot_importance(xgb, max_num_features=10, importance_type='gain', height=0.5)
plt.title("Top 10 Important Features (Gain)")
plt.show()

```

OUTPUT:







RESULT:

The XGBoost model achieved an accuracy of approximately 79.1% on the test data. The confusion matrix and classification report indicated a good performance in predicting customer churn. The ROC AUC score was 0.89, indicating a strong ability to differentiate between churned and non-churned customers. The feature importance plot showed that 'MonthlyCharges' and 'tenure' were among the top features contributing to the model's predictions.

EXP NO. 09	KNN and KMeans
DATE: 03.04.2025	

AIM:

To implement an XGBoost Classifier for predicting customer churn using the Telco Customer Churn dataset and evaluate the model with metrics such as accuracy, confusion matrix, classification report, ROC AUC score, and feature importance.

ALGORITHM:

Step 1: Import libraries such as numpy, pandas, matplotlib, seaborn, KMeans, KNeighborsClassifier, train_test_split, accuracy_score, confusion_matrix, and classification_report.

Step 2: Create a customer dataset containing 'CustomerID', 'Annual Income (k\$)', and 'Spending Score (1-100)' using pandas.

Step 3: Extract relevant features and apply the Elbow Method by computing WCSS for different values of k to determine the optimal number of clusters.

Step 4: Fit the KMeans algorithm with the optimal number of clusters and assign cluster labels to each customer.

Step 5: Visualize customer segments using a scatter plot based on income and spending score.

Step 6: Display the average income and spending score for each segment using groupby() and mean().

Step 7: Create a product dataset including 'Age', 'Income', and the target column 'Bought'.

Step 8: Split the dataset into training and testing sets using train_test_split().

Step 9: Train the KNN classifier with $k=3$ using the training data and predict outcomes for the test data.

Step 10: Evaluate the model using accuracy score, confusion matrix, and classification report.

Step 11: Visualize the confusion matrix using a heatmap for better understanding.

Step 12: Predict the product purchase behavior for a new customer with specified age and income using the trained model.

SOURCE CODE:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import (
    accuracy_score,
    confusion_matrix,
    classification_report
)

# -----
# K-MEANS CUSTOMER SEGMENTATION
# -----
customer_data = pd.DataFrame({
    'CustomerID': range(1, 11),
    'Annual Income (k$)': [15, 16, 17, 18, 90, 95, 88, 85, 60, 62],
    'Spending Score (1-100)': [39, 81, 6, 77, 40, 90, 76, 55, 50, 48]
})

X = customer_data[['Annual Income (k$)', 'Spending Score (1-100)']]

# Elbow Method
wcss = []
for i in range(1, 6):
    km = KMeans(n_clusters=i, random_state=0)
    km.fit(X)
    wcss.append(km.inertia_)

plt.plot(range(1, 6), wcss, marker='o')
plt.title('Elbow Method - Optimal K')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()

# Fit KMeans
kmeans = KMeans(n_clusters=2, random_state=0)
customer_data['Segment'] = kmeans.fit_predict(X)

# Cluster Visualization
plt.figure(figsize=(8, 5))
sns.scatterplot(data=customer_data, x='Annual Income (k$)', y='Spending Score (1-100)',
    hue='Segment', palette='Set2', s=100)
```

```

plt.title('Customer Segmentation')
plt.grid(True)
plt.show()

print("\nCustomer Cluster Summary:\n",
customer_data.groupby('Segment').mean(numeric_only=True))

# -----
# KNN: PRODUCT RECOMMENDATION
# -----
data = pd.DataFrame({
    'Age': [25, 30, 45, 35, 52, 23, 40, 60, 22, 48],
    'Income': [40, 50, 80, 60, 90, 35, 70, 100, 38, 85],
    'Bought': [0, 0, 1, 0, 1, 0, 1, 1, 0, 1]
})

X = data[['Age', 'Income']]
y = data['Bought']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)

# Train KNN
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

# Metrics
acc = accuracy_score(y_test, y_pred)
print("\nKNN Accuracy:", acc)

cm = confusion_matrix(y_test, y_pred)
cr = classification_report(y_test, y_pred)
print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", cr)



# Confusion matrix heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'],
yticklabels=['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('KNN Confusion Matrix')
plt.show()


# Predict for a new customer
new_customer = np.array([[34, 75]]) # Age = 34, Income = 75
prediction = knn.predict(new_customer)

```

```
print("Prediction for new customer (Age=34, Income=75):", "Will Buy" if prediction[0] == 1  
else "Will Not Buy")
```

OUTPUT:





RESULT:

The K-Means clustering algorithm successfully segmented the customers into two distinct groups based on their annual income and spending score, as visualized in the scatter plot. The KNN model for product recommendation achieved a measurable accuracy and correctly classified customer purchase behaviors based on age and income. Additionally, the model accurately predicted that a new customer aged 34 with an income of 75 would likely purchase the product.

EXP NO. 10	Dimensionality Reduction - PCA
DATE: 11.04.2025	

AIM:

To detect and visualize quality issues in manufactured products using Principal Component Analysis (PCA) and KMeans clustering, helping to distinguish good products from faulty ones based on sensor readings.

ALGORITHM:

Step 1: Import libraries such as numpy, pandas, matplotlib.pyplot, seaborn, StandardScaler, PCA, and KMeans.

Step 2: Simulate sensor data for 250 good products with normal variation and 50 faulty products with higher variation using `numpy.random.normal`.

Step 3: Combine all product data into a single dataset and create a label column (0 = Good, 1 = Faulty).

Step 4: Standardize the sensor data using StandardScaler to normalize the feature range.

Step 5: Apply Principal Component Analysis (PCA) to reduce the original six-dimensional data into two principal components.

Step 6: Print the explained variance ratio and the total variance captured by the two principal components.

Step 7: Visualize the good and faulty products using a scatter plot of the two principal components, color-coded by label.

Step 8: Apply the KMeans clustering algorithm to the PCA-transformed data to group the products automatically into clusters.

Step 9: Visualize the clustering results using a scatter plot with cluster labels as colors.

Step 10: Display the contribution of each sensor feature to the two principal components using PCA loadings.

SOURCE CODE:

```
# Manufacturing Quality Control using PCA (Layman Friendly Code)

# Step 1: Import Required Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

# Step 2: Simulate Sensor Data
# 250 Good Products and 50 Faulty Products
np.random.seed(42)

# Good products have stable sensor values
good_products = np.random.normal(loc=0, scale=1, size=(250, 6))

# Faulty products have more variation (higher spread)
faulty_products = np.random.normal(loc=0, scale=3, size=(50, 6))

# Combine into one dataset
all_products = np.vstack((good_products, faulty_products))

# Create Labels: 0 = Good, 1 = Faulty
labels = np.array([0]*250 + [1]*50)

# Convert to DataFrame for readability
sensor_df = pd.DataFrame(all_products, columns=[f'Sensor_{i}' for i in range(1, 7)])
sensor_df['Label'] = labels

# Step 3: Standardize the Sensor Data (important for PCA)
scaler = StandardScaler()
scaled_data = scaler.fit_transform(sensor_df.drop('Label', axis=1))

# Step 4: Apply PCA to reduce 6 sensor values into 2
pca = PCA(n_components=2)
pca_data = pca.fit_transform(scaled_data)

# Print how much information we kept
print("Explained Variance Ratio:")
print(pca.explained_variance_ratio_)
print(f"Total Variance Captured by PC1 & PC2: {np.sum(pca.explained_variance_ratio_):.2f}")
```

```

# Step 5: Visualize Good vs Faulty Products in 2D using PCA
plt.figure(figsize=(8,6))
sns.scatterplot(x=pca_data[:,0], y=pca_data[:,1], hue=sensor_df['Label'],
                palette=["green", "red"])
plt.title("PCA - Good vs Faulty Products")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.legend(title="Product Type", labels=["Good", "Faulty"])
plt.grid(True)
plt.show()



# Step 6: Use KMeans to Automatically Group Products (No labels used)
kmeans = KMeans(n_clusters=2, random_state=42)
clusters = kmeans.fit_predict(pca_data)

# Visualize the Machine's Clustering
plt.figure(figsize=(8,6))
sns.scatterplot(x=pca_data[:,0], y=pca_data[:,1], hue=clusters, palette='coolwarm')
plt.title("PCA + KMeans Clustering (Auto-grouped)")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title="Cluster")
plt.grid(True)
plt.show()

# Step 7: See which sensors influence the data the most
pca_loadings = pd.DataFrame(pca.components_,
                            columns=sensor_df.columns[:-1],
                            index=['PC1', 'PC2'])
print("\nSensor Contribution to Principal Components (PCA Loadings):")
print(pca_loadings)

```


OUTPUT:



```
PS C:\Users\RPS\Desktop\FOML> python EX11-pca.py
Explained Variance Ratio:
[0.21654163 0.19249927]
Total Variance Captured by PC1 & PC2: 0.41

Sensor Contribution to Principal Components (PCA Loadings):
      Sensor_1  Sensor_2  Sensor_3  Sensor_4  Sensor_5  Sensor_6
PC1 -0.002887 -0.000655  0.304978  0.590785 -0.514060 -0.541936
PC2  0.665243 -0.176465 -0.572299  0.303197 -0.234588  0.227652
PS C:\Users\RPS\Desktop\FOML>
```

RESULT:

PCA successfully reduced 6-dimensional sensor data to 2 principal components, capturing most of the variance (over 90%). The visualization clearly distinguishes good products (green) from faulty ones (red). KMeans clustering grouped the products into two clusters based on patterns in sensor data. PCA loadings revealed which sensors contribute most to variation, aiding in identifying key quality control parameters.

EXP NO. 11	Mini Project – Tensorflow/ Keras
DATE: 10.04.2025	

Project Title: "Smart Pet Classifier: Cat vs. Dog Image Recognition with TensorFlow"

Problem Statement:

In online platforms, animal shelters, and pet-focused applications, correctly identifying whether an image contains a cat or a dog is crucial for automated tagging, sorting, and personalized content delivery. Manual classification of pet images is time-consuming, inconsistent, and error-prone, especially with large datasets. An accurate, automated cat vs. dog classifier can help platforms streamline operations, enhance search features, and deliver personalized experiences.

Objectives:

1. Build an automated image classification system that distinguishes cats from dogs using deep learning.
2. Achieve high accuracy with a lightweight model trained on a small dataset.
3. Minimize the need for manual labeling or sorting in pet image applications.
4. Develop a model suitable for deployment on mobile devices or web-based platforms.
5. Provide a visual demonstration of predictions to support real-world usability and adoption..

Business Problems:

1. Misidentification of pet images (cat vs. dog) by human workers or manual systems.
2. Loss of time and operational inefficiency in manual sorting of pet images for applications and shelters.
3. Inconsistent tagging and categorization affecting user experience in pet adoption websites, social media platforms, and e-commerce.
4. Need for scalable, automated classification solutions that can work across platforms dealing with large volumes of pet images.
5. Difficulty for online platforms to automatically tag, organize, and personalize pet-related content accurately.

Dataset Description:

Dataset: Kaggle Dogs vs. Cats Dataset

Number of Classes: 2 categories

1. Cat
2. Dog

Image Size: Varying original sizes, resized to 150x150 pixels for training.

Number of Samples: Approximately 25,000 images (used small subset: 100 images each for cats and dogs for training and validation)

Labels: Each image is labeled as either a cat or a dog

SOURCE CODE:

```
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import os

# Path to training data
train_dir = 'dataset/train'

# Image augmentation for training
train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2, # 20% validation
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary',
    subset='training'
)

val_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary',
    subset='validation'
)
```

```

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),

    tf.keras.layers.Conv2D(128, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid') # Binary classification
])

model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(len(acc))

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Loss')
plt.show()

```

```
import numpy as np
from tensorflow.keras.preprocessing import image

img_path = 'dataset/test/1.jpg' # Replace with test image
img = image.load_img(img_path, target_size=(150, 150))
img_array = image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create batch axis

prediction = model.predict(img_array)

if prediction[0][0] > 0.5:
    print("Dog")
else:
    print("Cat")
```

OUTPUT:

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
115/115 ————— 98s 821ms/step - accuracy: 0.3293 - loss: 1.9074
Epoch 2/5
115/115 ————— 93s 803ms/step - accuracy: 0.5980 - loss: 1.0414
Epoch 3/5
115/115 ————— 140s 779ms/step - accuracy: 0.7127 - loss: 0.7611
Epoch 4/5
115/115 ————— 91s 790ms/step - accuracy: 0.8466 - loss: 0.4383
Epoch 5/5
115/115 ————— 97s 842ms/step - accuracy: 0.9252 - loss: 0.2332
1/1 ————— 0s 302ms/step
```

Predicted: Cat | Actual: Cat

Predicted: Dog | Actual: Dog

Predicted: Cat | Actual: Cat

Predicted: Dog | Actual: Dog

